Brain Tumour Magnetic Resonance Image Classification with Perceptual-Hash Deduplication: A Comparative Study Using Advanced Deep Learning Models

Parth Mital[a], Ved Lawand[a], Biji C. L.[a,*], Muzaffar Ahamad Dar[b] and George Priya Doss C.[b]

[a]School of Computer Science and Engineering, Vellore Institute of Technology, Vellore, Tamil Nadu, India
[b]School of Bio Sciences and Technology, Vellore Institute of Technology, Vellore, Tamil Nadu, India

ARTICLE INFO

Keywords:
Brain tumour MRI
deep learning
perceptual hashing
shortcut learning
vision transformers

ABSTRACT

Magnetic resonance imaging (MRI) is the standard investigation for brain tumours, but published deep learning accuracies are difficult to compare because studies differ in data, splitting and evaluation, and near-duplicate images from shared public sources leak between training and test partitions. The contribution to artificial intelligence is a controlled evaluation protocol: label-aware perceptual-hash deduplication before splitting, a group-aware split into five cross-validation folds and a held-out test set, an equal number of optimiser steps for every model, temperature scaling, fold ensembles with bootstrap confidence intervals and Holm-corrected McNemar tests, and attention and background-only shortcut audits. The engineering application is the selection of a deployable tumour detector and subtype classifier. Nine ImageNet-pretrained convolutional, transformer and hybrid backbones were benchmarked on binary detection and three-way subtype classification. Deduplication reduced the binary detection pool from 3,253 to 1,185 images and the subtype classification pool from 5,828 to 4,730, and curation ran at about 85 images per second on four CPU cores. On 171 binary detection test images, CoAtNet-0, EfficientNetB0 and ViT-B/16 reached 99.42% accuracy (95% confidence interval 98.25–100%), and no pair of models differed significantly. The subtype classification grouping left only 45 test images, on which six models were error-free, so they cannot rank the models. Models retrained with the head masked out still exceeded the majority-class rate for 7 of 9 binary detection and all 9 subtype classification backbones, reaching 83.04% and 97.78% accuracy, so these public collections carry label information outside the brain. EfficientNetB0 matched the binary detection leaders with 4.01 million parameters.

Corresponding author

Biji C. L.
School of Computer Science and Engineering, Vellore Institute of Technology, Vellore 632014, Tamil Nadu, India
E-mail: biji.cl@vit.ac.in
Phone: +91 94461 71877

*Corresponding author
biji.cl@vit.ac.in (B.C. L.)